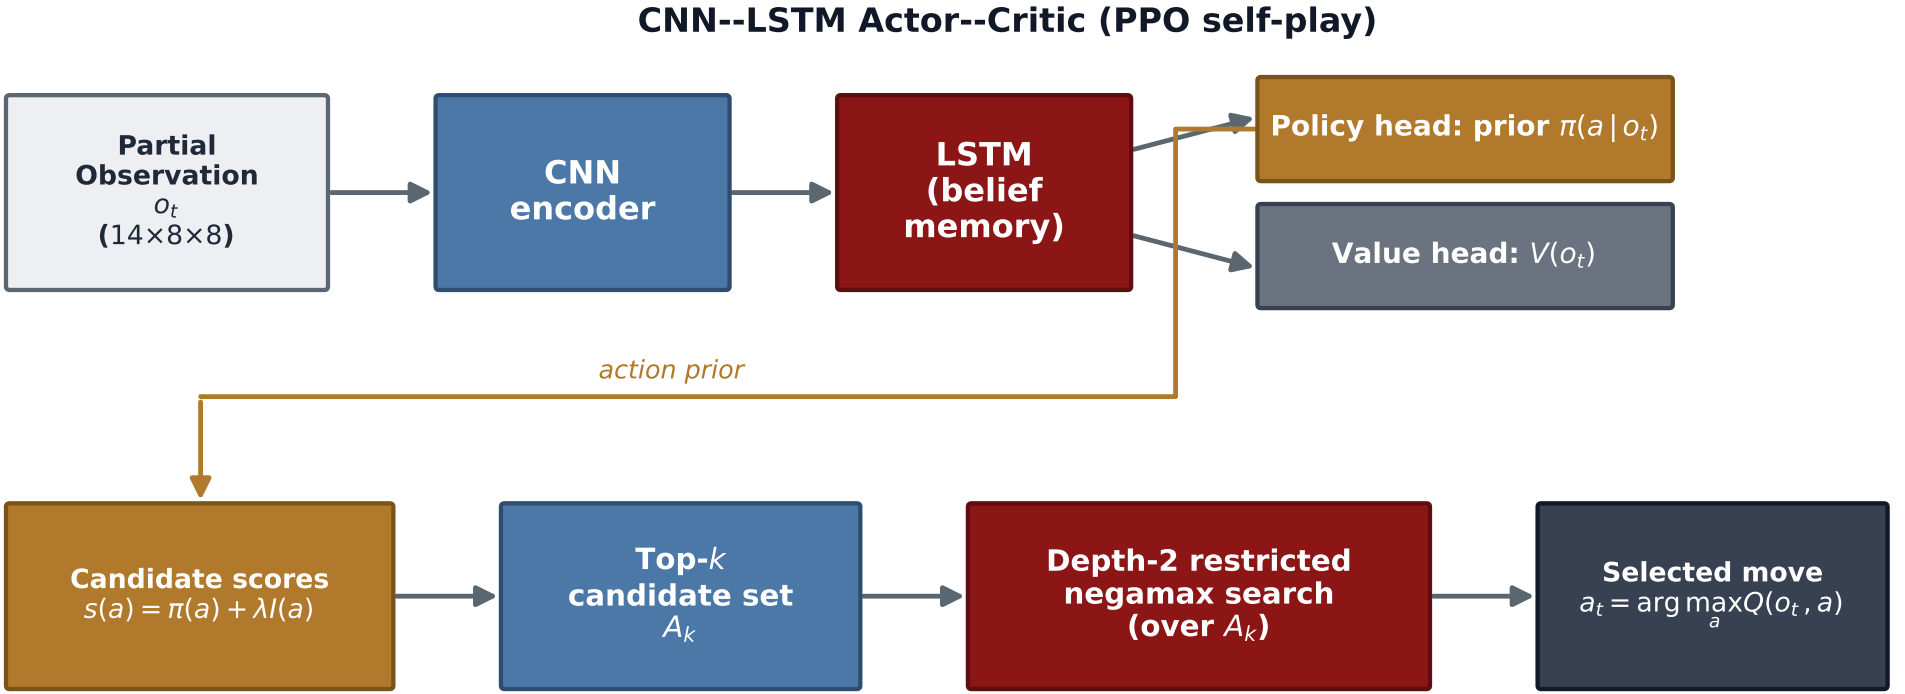
CNN--LSTM Actor--Critic (PPO self-play)
Partial Observation o_t (14×8×8)
CNN encoder
LSTM (belief memory)
Policy head: prior $\pi(a \mid o_t)$
Value head: $V(o_t)$
action prior
Candidate scores $s(a) = \pi(a) + \lambda I(a)$
Top-k candidate set A_k
Depth-2 restricted negamax search (over A_k)
Selected move $a_t = \arg\max_a Q(o_t, a)$
Inference: prior-guided candidate selection → restricted search